

Abstract

Large language models produce confident but factually incorrect outputs at rates that span three to ninety-one percent across deployment domains, a failure mode that increased parameter count alone does not close. This thesis proposes Confidence-Aware Episodic Memory with Self-Improvement (CAEM), an architectural intervention that treats hallucination reduction as a system-level problem rather than a post-hoc detection problem. A single query passes through a verified episodic memory store, a confidence-aware three-tier router with an independent safety override and a retrieval-utility classifier, and a nine-signal verifier ensemble spanning internal calibration, sample-set agreement, external grounding, and question-answer relevance, whose per-benchmark calibrated-probability composite gates storage at two fixed cut-points held constant across cycles. Self-improvement is delivered through low-rank adaptation over a frozen Qwen-2.5-3B-Instruct backbone, with every cycle boundary refitting calibration on a fresh fold under a multi-modal retention guard whose asymmetric rollback reverts the adapter while preserving memory. Against the matched zero-shot baseline on the same backbone, on a five-benchmark panel split into training (FEVER, TriviaQA, CommonsenseQA) and transfer (TruthfulQA, StrategyQA), the within-trajectory best cycle reaches a pooled exact-match gain of +20.3% and a pooled composite hallucination metric reduction of −29.4% on the training panel; against the locked external eight-baseline panel, CAEM leads every baseline on pooled composite hallucination at both cycles and on pooled exact match at the best cycle, reaching exact-match parity with the strongest baselines at trajectory end.


Research Objective

- Verified episodic memory with a four-outcome decision** whose admission rule is finer than a single binary threshold.
- Pre-generation confidence, three-tier routing with safety override, and a retrieval-utility classifier** computed from a single short forward pass and calibrated via per-benchmark temperature scaling on a disjoint fold.
- Nine-signal post-generation verifier with a calibrated probability composite** across four families: internal calibration, sample-set agreement, external grounding under three aggregation rules, and question-answer relevance.
- Fixed-threshold storage gate with two registered cut-points** held constant across cycles and one abstention floor.
- Retroactive review and deferred reconsideration at cycle boundaries** that rescore stored episodes under the updated model and promote deferred entries past threshold.
- Constrained self-improvement with a parameter-efficient adapter and a multi-modal retention guard** that aborts the cycle with an adapter rollback when the worst probe ratio falls below a fixed floor.
- Empirical validation against eight external baselines and pre-registered ablations** on a five-benchmark panel under matched protocol.

Methodology

The architecture organises a single query into an eight-stage pipeline traversed once per query and wrapped in a per-cycle update loop. Stage 1 produces a pre-routing confidence from a single short forward pass. Stage 2 encodes the query and probes a similarity-indexed episodic memory for the nearest stored entry. Stage 3 combines the similarity and stored-quality scores into a dispatch score and routes to one of three tiers, with a safety override on low confidence. Stage 3b, a retrieval-utility classifier, refines the dispatch into zero-shot or retrieval-augmented tier per query. Stage 4 executes the chosen tier; Stage 5 passes the answer through a nine-signal post-generation verifier; Stage 6 applies a four-outcome decision (store, defer, abstain, discard); Stage 7 returns the user-facing action. Stage 8, at cycle boundaries, runs retroactive re-verification, deferred reconsideration, memory consolidation, and a regularised fine-tune under a retention guard with asymmetric rollback: weights are rolled back on

Model Architecture



The eight-stage pipeline routes each query through pre-routing confidence, memory lookup, three-tier dispatch with the Stage 3b retrieval-utility classifier ($\tau_{RUC} = 0.375$), tier execution, nine-signal verification, the four-outcome decision under fixed cut-points $\tau_{store} = 0.60$ / $\tau_{defer} = 0.45$, and output; Stage 8 closes every cycle boundary with retroverify ($\tau_{retrain} = 0.50$), deferred reconsider (TTL = 2), consolidate, LoRA fine-tune ($r = 32$, $\alpha = 64$, $\tau_{train} = 0.70$), and the multi-modal retention probe under floor $\rho_{min} = 0.93$ with asymmetric rollback (adapter reverted, memory preserved). Backbone Qwen-2.5-3B-Instruct, frozen across the cycle.

Per-Cycle Trajectory

		Pooled headline Retention probes (ratio vs. pristine)							
Configuration	Cycle n /bench	EM	CHM	MMLU	TQA-test	CSQA-test		Outcome	
<i>Reference floor (no CAEM architecture; same Qwen-2.5-3B-Instruct backbone)</i>									
B1 (zero-shot)	—	300	0.507	0.120	—	—	—	reference	
<i>CAEM trajectory (full architecture: memory + verifier + SIL + retention guard)</i>									
CAEM	C0	300	0.541	0.135	—	—	—	baseline	
CAEM	C1	300	0.543	0.112	1.041	1.013	1.006	COMMIT	
CAEM	C2	300	0.544	0.107	0.975	<u>0.949</u>	0.982	COMMIT ★	
CAEM	C3	300	0.532	0.115	0.992	0.987	0.976	COMMIT	
CAEM	C4	300	0.532	0.115	1.018	0.886	0.982	ABORT	
CAEM	C5	300	0.517	0.109	0.999	1.013	<u>0.958</u>	COMMIT	
Headline contrast (CAEM C2 vs. B1 zero-shot): $\Delta\text{EM} = +0.037$ (+7.4% relative); $\Delta\text{CHM} = -0.013$ (−10.8% relative).									

Headline Result

Against the matched zero-shot baseline B1 on the same backbone, the training-panel architectural-contribution slice at the within-trajectory best cycle (C2) reaches a pooled exact-match gain of +20.3% relative and a pooled composite hallucination metric reduction of −29.4% relative; the trajectory-end cycle holds at +13.7% EM and −20.3% CHM. CAEM beats every external baseline on composite hallucination at both reference cycles and on exact match at the best cycle, reaching exact-match parity with the strongest baselines at trajectory end.

$$\Delta\text{EM} = +20.3\% \text{ rel.}$$

$$\Delta\text{CHM} = -29.4\% \text{ rel.}$$

(training panel, CAEM C2 vs. B1 zero-shot baseline)

External Baseline Comparison

# System	EM	CHM	vs. CAEM C2 (best)		vs. CAEM C5 (final)	
			ΔEM	ΔCHM	ΔEM	ΔCHM
CAEM C2 (best)	0.544	0.107	—	—	+0.027	−0.002
CAEM C5 (final)	0.517	0.109	−0.027	+0.002	—	—
B1 zero-shot	0.507	0.120	+0.037 (+7.4%)	−0.013 (−11.0%)	+0.010 (+2.0%)	−0.011 (−9.3%)
B2 chain-of-thought	0.500	0.123	+0.044 (+8.8%)	−0.016 (−12.7%)	+0.017 (+3.4%)	−0.014 (−11.0%)
B3 RAG	0.458	0.148	+0.086 (+18.8%)	−0.041 (−27.7%)	+0.059 (+13.0%)	−0.039 (−26.4%)
B4 CoT + RAG	0.464	0.142	+0.080 (+17.2%)	−0.035 (−24.6%)	+0.053 (+11.5%)	−0.033 (−23.2%)
B5 five-shot CoT	0.405	0.136	+0.139 (+34.2%)	−0.029 (−21.4%)	+0.112 (+27.6%)	−0.027 (−19.9%)
B6 vanilla fine-tune (C5)	0.490	0.127	+0.054 (+11.0%)	−0.020 (−15.9%)	+0.027 (+5.6%)	−0.018 (−14.3%)
B7 FLARE	0.525	0.173	+0.019 (+3.5%)	−0.066 (−38.0%)	−0.008 (−1.5%)	−0.064 (−36.8%)
B8 semantic entropy	0.505	0.120	+0.039 (+7.8%)	−0.013 (−10.5%)	+0.013 (+2.5%)	−0.011 (−8.9%)
Wins-against count (CAEM beats baseline on ΔEM and ΔCHM jointly): CAEM C2 beats 8/8; CAEM C5 beats 7/8.						

Retrieval-Utility Classifier Ablation

Benchmark	Exact match			Composite hallucination			Tier 3 share		
	Off	On	Δ (pp)	Off	On	Δ (% rel)	Off	On	Δ (pp)
<i>Training panel</i>									
FEVER	0.503	0.503	0.0	0.117	0.120	+2.5%	92.3%	70.7%	−21.7
TriviaQA	0.437	0.463	+2.7	0.123	0.121	−2.0%	95.7%	68.0%	−27.7
CommonsenseQA	0.603	0.743	+14.0	0.131	0.086	−34.4%	57.7%	20.3%	−37.3

<i>Transfer panel</i>									
StrategyQA	0.617	0.640	+2.3	0.119	0.128	+8.1%	100.0%	76.7%	−23.3
TruthfulQA	0.210	0.310	+10.0	0.116	0.119	+2.9%	99.0%	40.0%	−59.0
Pooled (all five)	0.474	0.532	+5.8	0.121	0.115	−5.2%	88.9%	55.1%	−33.8

References

- Huang et al. A survey on hallucination in large language models. *ACM TOIS* 2024.
- Min et al. FactScore: fine-grained atomic evaluation of factual precision in long-form text generation. *EMNLP* 2023.
- Guo et al. On calibration of modern neural networks. *ICML* 2017.
- Lin et al. TruthfulQA: measuring how models mimic human falsehoods. *ACL* 2022.
- Lewis et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *NeurIPS* 2020.
- Wang et al. Self-consistency improves chain-of-thought reasoning. *ICLR* 2023.
- Farquhar et al. Detecting hallucinations in large language models using semantic entropy. *Nature* 2024.
- Hu et al. LoRA: low-rank adaptation of large language models. *ICLR* 2022.
- Tang et al. MiniCheck: efficient fact-checking on grounding documents. *ACL* 2024.
- Cherian et al. Enhanced conformal prediction with boosted logistic regression. *NeurIPS* 2024.